

Spark the Definitive Guide 2nd Edition

Chapter 04

Structured API Overview

Structured API Overview

Text Book



Objectives and Outcomes

- ▶ Introduced to Spark's Structured APIs, Datasets, DataFrames, and SQL Views
- ▶ Learn how Spark transforms into a physical execution plan on a cluster

Review

So far:

- ▶ We learned about Spark's programming model
- ▶ We learned how to run production code
- ▶ We were introduced to type-safe data structures in Spark
- ▶ We were introduced to Structured Streaming on Spark
- ▶ We were introduced to Machine Learning on Spark
- ▶ We were introduced to 3rd party Spark packages

API Overview 66

- ▶ Three datatypes in Spark:
 - ▶ DataFrames
 - ▶ Can you define this term?
 - ▶ Datasets
 - ▶ Can you define this term?
 - ▶ SQL Tables and Views
 - ▶ Can you define these terms?
- ▶ With these data types we can manipulate disparate types of data
 - ▶ Unstructured log files
 - ▶ Semi-structured CSV files
 - ▶ Structured Parquet files

Structured API concepts

- ▶ These concepts refer to both *batch* and *streaming*
 - ▶ Code should easily switch between the two
 - ▶ We will cover Streaming later in the course, Chapter 20

Structured Collections

- ▶ Spark has two notions of structured collections:
 - ▶ Datasets and DataFrames
- ▶ Each are distributed table-like collections with well defined rows and columns
 - ▶ Each row must have the same number of columns
 - ▶ Both are **immutable**
 - ▶ Both allow for lazily evaluated plans that are only deployed when an **action** is called

Schemas

- ▶ A **schema** defines the column names and data types of the column
 - ▶ Schemas can be defined manually or inferred
 - ▶ Schema on Read
- ▶ All of Spark actions take place in the internal Spark language called Catalyst
 - ▶ We don't write in this language but the JVM allows us to write in higher level languages that convert to Catalyst

DataFrames vs Datasets

- ▶ DataFrames have types of a sort. . .
 - ▶ These are maintained by Spark internally
 - ▶ Schema only checked at *runtime*
- ▶ Datasets are typed DataFrames
 - ▶ Only available in Scala and Java
 - ▶ Enforce type at compile time
 - ▶ P. 54

Overview of Structured Spark Types

- ▶ Spark is effectively a programming language of its own
 - ▶ Uses the *Catalyst* engine internally to maintain type information
- ▶ This code does not do math in Scala, but Catalyst:
 - ▶ `scala val df = spark.range(500).toDF("number")
df.select(df.col("number") + 10)`

DataFrames vs. Datasets

- ▶ DataFrame schema checked at *runtime*
- ▶ Dataset schema checked at *compile time*
 - ▶ Datasets only available in Java and Scala
 - ▶ Why?
- ▶ DataFrames are Datasets of type Row
 - ▶ Type Row is Spark's internal optimized in-memory format for computation P.54
- ▶ Even without Datasets in Python and R, we are still always working on an optimized in-memory datatype

Columns and Rows

- ▶ Columns represent a 3 types of data:
 - ▶ A *simple type* like an integer or string
 - ▶ A *complex type* like an array or map
 - ▶ A *null value* :!
- ▶ A row is nothing more than a record of data
- ▶ Each record in a DataFrame must be of type Row
- ▶ Rows can be created in numerous ways:
 - ▶ Via SQL statements
 - ▶ DataSources (ingesting)
 - ▶ dynamically and in memory
 - ▶ `spark.range(2).toDF().collect()`

Spark Types

- ▶ You can import the types library you want to work with in Scala
 - ▶ `import org.apache.spark.sql.types._`
 - ▶ `val b = ByteType`
- ▶ You can import the types library you want to work with in Java
 - ▶ `import org.apache.spark.sql.types.DataTypes;`
 - ▶ `ByteType x = DataTypes.ByteType;`
- ▶ You can import the types library you want to work with in Python
 - ▶ `from pyspark.sql.types import *`
 - ▶ `b = ByteType()`
- ▶ Page 56 has an entire table of all the data type libraries available

Overview of Structured API Execution

- ▶ Structured API execution happens in 4 steps on Page 58:
 - ▶ Write your DataFrame/Dataset/SQL code
 - ▶ If valid code, Spark converts this to a *Logical Plan*
 - ▶ Spark transforms this *Logical Plan* to a *Physical Plan*, checking for optimizations along the way
 - ▶ Spark then executes this *Physical Plan* on the cluster




Figure 4-1. The Catalyst Optimizer

Logical and Physical Planning

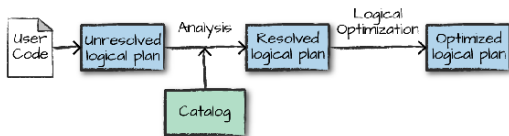


Figure 4-2. The structured API logical planning process

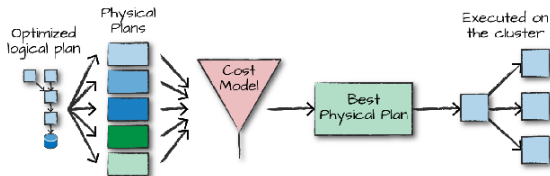


Figure 4-3. The physical planning process

Logical Planning

- ▶ The first phase takes the user code and converts it into a logical plan
 - ▶ Purely to convert the code into the most optimized version
 - ▶ Spark has an unresolved logical plan
 - ▶ Your code may compile, but what if the table name or column name is wrong?
 - ▶ Spark uses a *catalog* - an internal repo of all table and DataFrame information
 - ▶ Then *resolves* column and *tables* in the *analyzer*
- ▶ The analyzer might reject an *unresolved logical plan*, otherwise pass it to the *Catalyst Optimizer*
 - ▶ A collection of rules that attempts to optimize the logical plan by pushing predicates or selections down

Physical Planning

- ▶ After an optimized plan is generated.
 - ▶ Spark begins to specify how this plan will be executed on the cluster
 - ▶ Creates multiple strategies and compares them via a *cost model*

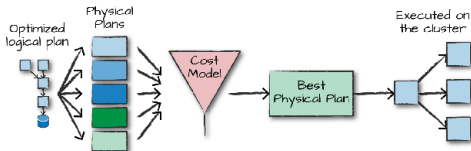


Figure 4-3. The physical planning process

- ▶ At execution time, Java bytecode is generated and the final result returned to the user

Conclusion

- ▶ We were introduced to Spark's Structured APIs, Datasets, DataFrames, and SQL Views
- ▶ We learned how Spark transforms a logical plan into a physical execution plan on a cluster

Questions

- ▶ Any questions?
- ▶ Read Chapter 05 and do any exercises in the book.