

Spark the Definitive Guide 2nd Edition

Chapter 05

Basic Structured Operations

Basic Structured Operations

Text Book



Objectives and Outcomes

- ▶ Introduce the tools we will use to manipulate DataFrames
- ▶ Focus on fundamental DataFrame operations
 - ▶ Understand what an expression is
 - ▶ Understand the difference between `Select` and `SelectExpr`
 - ▶ Understand how to add columns and rows to a DataFrame
 - ▶ Understand how to take random samples from DataFrames

Review

So far:

- ▶ We were introduced to Spark's Structured APIs, Datasets, DataFrames, and SQL Views
- ▶ We learned how Spark transforms a logical plan into a physical execution plan on a cluster
- ▶ Learned how DataFrames consist of a series of records
- ▶ Learned how DataFrames are of type Row and have a number of columns
- ▶ Learned that schemas define the name and type of data in each column
- ▶ Learned that Partitioning of the DataFrame defines the layout of the DataFrames physical distribution on the cluster

Create a DataFrame

```
df = spark.read.format("json").load("Spark-The-Definitive-Guide")  
  
# Let's take a look at the defined schema  
df.printSchema()
```

- ▶ JSON

- ▶ is a lightweight, text-based data interchange format.

Schemas

- ▶ Schemas tie everything together
- ▶ Schema defines the column names and column types of a DataFrame
 - ▶ Schema can be applied on read or inferred or declared
- ▶ For Ad-hoc data usually *schema-on-read* is good enough
 - ▶ Though it can be a bit slow when dealing with text-based file formats like:
 - ▶ CSV
 - ▶ JSON
- ▶ Schema-on-read can lead to precision problems
 - ▶ If a column is really of type LONG but the numbers are smaller and interpreted as type INT
- ▶ Spark can be used for **ETL**:
 - ▶ Extraction
 - ▶ Transform
 - ▶ Load In these cases it is best to provide the schema to ensure type matches

JSON Object

```
spark.read.format("json").load("Spark-The-Definitive-Guide/
```

```
# This datatype is returned from the previous command  
# StructType(List(StructField  
# (DEST_COUNTRY_NAME,StringType,true),  
# StructField  
# (ORIGIN_COUNTRY_NAME,StringType,true),  
# StructField(count,LongType,true)))
```

- ▶ A schema is a StructType made up of a number of fields
 - ▶ StructFields have a name, type, and a Boolean flag indicating if they take nulls
 - ▶ If types in the data at run-time do not match the schema, Spark will throw an error

Declare a Schema

```
from pyspark.sql.types import StructField, StructType, StringType

myManualSchema = StructType([StructField("DEST_COUNTRY_NAME", StringType(), True),
                               StructField("ORIGIN_COUNTRY_NAME", StringType(), True)])

df = spark.read.format("json").schema(myManualSchema)
    .load("Spark-The-Definitive-Guide/data/flight-data/json/2014")
```

Columns and Expressions

- ▶ Columns can be selected, manipulated, and removed from DataFrames
 - ▶ These operations are referred to as *expressions*
 - ▶ Must use Spark to manipulate Rows (logical collection of Rows is a column)
 - ▶ Must be in the context of a DataFrame
 - ▶ To work on columns use the *col* or *column* functions
 - ▶ We will stick to using the *col* function
 - ▶ Columns are not resolved until compared to the *catalog* at run-time
 - ▶ Column and table resolution happen in the *analyzer* phase

```
from pyspark.sql.functions import col, column
```

```
col("someColumnName")
```

```
column("someColumnName")
```

Column Reference

- ▶ If you need to explicitly reference a column you can
- ▶ Think of it as a namespace way to reference columns in different DataFrames that have the same name
 - ▶ `df.col("count")`

Columns as Expressions

- ▶ What is an *expression*?
 - ▶ A set of *transformations* on one or more values in a *record* in a DataFrame
- ▶ You can use a `col()` and perform a transformation on a column
- ▶ You can use an `expr()` to parse transformations and column references
 - ▶ These references can subsequently be passed into further transformations
 - ▶ `expr("someCol - 5")` and `col("someCol") - 5` and `expr("someCol") - 5` all evaluate the same
 - ▶ Spark compiles these to the same logic tree
- ▶ Columns are just expressions
- ▶ Columns and transformations of those columns compile to the same logical plan
 - ▶ `((col("someCol") + 5) * 200) - 6` < `col("otherCol")`

Directed Acyclic Graph

- ▶ This is also represented by in Python (64):



```
from pyspark.sql.functions import expr expr("((someCol
```

- ▶ Previous expression is actually valid SQL code

- ▶ This means you can write your expressions as DataFrame code or as SQL expressions and get the same performance characteristics

Accessing a DataFrames Columns

- ▶ How can you see a DataFrame's columns?

```
spark.read.format("json").load("The-Definitive-Guide-To
```

Records and Rows 65

- ▶ Review: Each row in a DataFrame is a single record
 - ▶ Represented as an object of type Row
- ▶ How to read the first row of a DataFrame:
 - ▶ `df.first()`
- ▶ Only DataFrames have schemas, Rows do not have a schema
- ▶ To create a *Row* you must append values in the correct “schema”
 - ▶ `from pyspark.sql import Row`
 - ▶ `myRow = Row("Hello", None, 1, False)`
- ▶ To access Rows, Python and R will autodetect the datatype
 - ▶ `myRow2`
 - ▶ `myRow0`
- ▶ Scala and Java will require casting or coercing the values
 - ▶ `myRow(0).asInstanceOf[String] // String`
 - ▶ `myRow.getInt(2)`

DataFrame Transformations

- ▶ When working with individual DataFrames:
 - ▶ We can add rows or columns
 - ▶ We can remove rows or columns
 - ▶ We can transform a row into a column
 - ▶ We can change the order of rows based on the values of columns




Figure 5-2. Different kinds of transformations

Creating DataFrames

- ▶ We can create DataFrames from raw sources
 - ▶ Chapter 9 will cover this in more detail
 - ▶ We can register raw data as a temporary view
 - ▶ Query it with SQL
- ▶ We can create a DataFrame on the fly by taking a set of rows and converting them to a DataFrame

Code Example 65

```
df = spark.read.format("json").load("data/flight-data/json/dfTable")
df.createOrReplaceTempView("dfTable")
```

```
import org.apache.spark.sql.Row
import org.apache.spark.sql.types.{StructField, StructType,
```

```
val myManualSchema = new StructType(Array(new StructField("
val myRows = Seq(Row("Hello", null, 1L))
val myRDD = spark.sparkContext.parallelize(myRows)
val myDf = spark.createDataFrame(myRDD, myManualSchema)
myDf.show()
```

```
// use can map Scala Seq directly to DataFrames, but Seq d
val myDf = Seq(("Hello",2,1L)).toDf("col1","col2","col3")
```

```
from pyspark.sql import Row
from pyspark.sql.types import StructField, StructType, Str
myManualSchema = StructType([
    StructField("some", StringType(), True),
    StructField("col", StringType(), True)
```

Select and selectExpr

- ▶ Use the *select* method when working with columns or expressions
- ▶ Use the *selectExpr* method when working with expressions in **strings**
- ▶ Both are found in `org.apache.spark.sql.functions`
 - ▶ `select` and `selectExpr` allow you to execute SQL queries on a `DataFrame`
 - ▶ `df.select("DEST_COUNTRY_NAME").show(2)`
 - ▶ **`SELECT`** `DEST_COUNTRY_NAME` **`FROM`** `dfTable` **`LIMIT`** `2`
- ▶ You can select multiple columns by using a comma

```
from pyspark.sql.functions import expr, col, column
df.select(
    expr("DEST_COUNTRY_NAME"),
    col("DEST_COUNTRY_NAME"),
    column("DEST_COUNTRY_NAME")
).show(2)
```

selectExpr

- ▶ If you find yourself typing a bunch of *select* then *expr* statements:
 - ▶ Then selectExpr is the convenient interface you want
 - ▶ We can add new columns to a DataFrame
- ▶ We can use selectExpr to build up complex expressions and create new DataFrames
 - ▶

```
df.selectExpr("*", ("DEST_COUNTRY_NAME = ORIGIN_COUNTRY_NAME"))
```
 - ▶

```
SELECT *, (DEST_COUNTRY_NAME = ORIGIN_COUNTRY_NAME) as w
```
- ▶ We can specify aggregations over an entire DataFrame
 - ▶

```
df.selectExpr("avg(count)", "count(distinct(DEST_COUNTRY_NAME))")
```
 - ▶

```
SELECT avg(count), count(distinct(DEST_COUNTRY_NAME)) FROM
```

Spark Literals

- ▶ Sometimes we need to pass a literal value, such as a constant
 - ▶ `from pyspark.sql.functions import lit`
 - ▶
`df.selectExpr(expr("*"), lit(1).alias("One")).show(2)`
 - ▶ This will come up when you need to check a Row value against a predetermined constant value
- ▶ Adding additional columns is possible: `.withColumn()`
 - ▶
`df.withColumn("withinCountry", expr("ORIGIN_COUNTRY_NAME"))`
 - ▶ This creates a column with a Boolean if the ORIGIN and DEST Country name match.
 - ▶ This can save much time in a lookup later on as you will not have to do String comparison
- ▶ Columns can be dropped as well
 - ▶ `df.drop("ORIGIN_COUNTRY_NAME").columns`
- ▶ You can cast columns as well
 - ▶ `df.withColumn("count2", col("count").cast("long"))`
 - ▶ Renaming a column is possible using the `.withColumnRenamed("existingColumnName", "newColumnName")`

Filter and Where Clauses 72

- ▶ In working with Spark DataFrames, you can use both `where` and `filter` on a DataFrame
 - ▶ `df.filter(col("count") < 2).show(2)`
 - ▶ `df.where("count < 2").show(2)`
 - ▶ More details in Chapter 11
- ▶ Both statements have the same output, `where` is a familiar term from SQL so the book will use that
- ▶ You can chain multiple `where` statements together, Spark will handle the expressions at run time
 - ▶

```
df.where(col("count") < 2).where(col("ORIGIN_COUNTRY_NAME"
```
 - ▶

```
SELECT * FROM dfTable WHERE count < 2 AND ORIGIN_COUNTRY
```
- ▶ You can access distinct results as we saw earlier in the chapter
 - ▶

```
df.select("ORIGIN_COUNTRY_NAME", "DEST_COUNTRY_NAME").di
```

Random Samples and Splits

- ▶ Sometimes you want to select a random sample of data for running a test on a small representative set
 - ▶ You can use the `sample` method on a `DataFrame`
 - ▶ `seed = 5`
`withReplacement = False`
`fraction = 0.5`
`df.sample(withReplacement, fraction, seed).count()`
- ▶ You can split a `DataFrame`
 - ▶ The seed definition is how the random selection is begun
 - ▶ `dataFrames = df.randomSplit([0.25, 0.75], seed)`
`dataFrames[0].count() > dataFrames[1].count()`

Concatenating and Appending Rows (Union)

- ▶ Previously we learned that DataFrames are **immutable**
 - ▶ How then can we append to a DataFrame?
 - ▶ In order to append to a DataFrame, you must **union** the original DataFrame along with the new DataFrame
 - ▶ Both DataFrames need to have the same schema and number of columns, otherwise the operation fails
 - ▶ The new DataFrame union can be assigned to a view or registered as a table to allow your code to reference it

```
from pyspark.sql import Row
schema = df.schema
newRows = [
    Row("New Country", "Other Country", 5L),
    Row("New Country 2", "Other Country 3", 1L)
]
parallelizedRows = spark.sparkContext.parallelize(newRows)
newDF = spark.createDataFrame(parallelizedRows, schema)

df.union(newDF).where("count = 1").where(col("ORIGIN_COUNTRY")
```


Sorting Rows

- ▶ You can sort the output of a DataFrame
 - ▶ You can use the sort and orderBy, they work exactly the same
 - ▶ They accept both column expressions and strings as well as multiple columns
 - ▶ Default is to sort ascending
 - ▶ You can sort nulls, asc_nulls_first, desc_nulls_first, asc_nulls_last, desc_nulls_last

```
df.sort("count").show(5)
```

```
df.orderBy("count", "DEST_COUNTRY_NAME").show(5)
```

```
df.orderBy(col("count"), col("DEST_COUNTRY_NAME")).show(5)
```

```
from pyspark.sql.functions import desc, asc
```

```
df.orderBy(expr("count desc")).show(2)
```

```
df.orderBy(col("count").desc(), col("DEST_COUNTRY_NAME").asc())
```

```
SELECT * FROM dfTable ORDER BY count DESC, DEST_COUNTRY_NAME
```

Repartition and Coalesce

- ▶ An important optimization opportunity is to partition the data according to a frequently filtered column
 - ▶ This controls the physical layout of data across the cluster
 - ▶ Including the partitioning scheme and the number of partitions
- ▶ Repartitioning will incur a full shuffle
 - ▶ This means you should only repartition when future number of partitions is greater than your current number of partitions
 - ▶ Or when you are looking to partition by a set of columns
 - ▶ `df.rdd.getNumPartitions()`
 - ▶ `df.repartitions(5)`
- ▶ If you know you will be filtering by a certain column often, you can repartition based on that column:
 - ▶ `df.repartition(col("DEST_COUNTRY_NAME"))`
 - ▶ You can specify how many partitions you want
 - ▶ `df.repartition(5, col("DEST_COUNTRY_NAME"))`

Conclusion

- ▶ This chapter covered basic DataFrame operations.
 - ▶ We learned the simple concepts and tools that you will need to be successful with Spark DataFrames
 - ▶ We learned what an expression is
 - ▶ We learned the difference between Select and SelectExpr
 - ▶ We learned how to add columns and rows to a DataFrame
 - ▶ We learned how to take random samples from DataFrames

Questions

- ▶ Any questions?
- ▶ Read Chapter 06 and do any exercises in the book.